

Power consumption prediction using Neural Networks, Decision Trees, Forest & Boosting Algorithms

MECE 4520 Final Project: Data Science for Mechanical Systems

Columbia University

By

Ali Danish, Avikal Sagar, Pei-Yu Chang, Peter Wu, Qixiao Zhang, Shevi Wolvovsky

Introduction

1. Power Consumption (PC) plays a critical role in a global economy due to the imbalance between energy production and demand.
2. Machine Learning and Deep Learning models have been widely used in PC predictions to help energy managers to control power systems better and improve energy usage.
3. In our project, we used **LSTMs (Neural Networks)**, **Decision Trees** and **Random Forests** and **Boosting algorithms** in order to achieve the same.
4. Power consumption from 3 power zones/stations were considered in **Tetouan (Morocco)**

Data Preprocessing

1. Imported the dataset from UCI Machine Learning repository [[Link](#)]
2. Converted and broke down 'Datetime' column (time series index) into - hours, minutes, quarter, day of week, day of month, day of year and month (time series features).
3. Plotted correlation matrix, boxplots and scatterplots (data visualization) to understand feature importance.


Our initial dataset



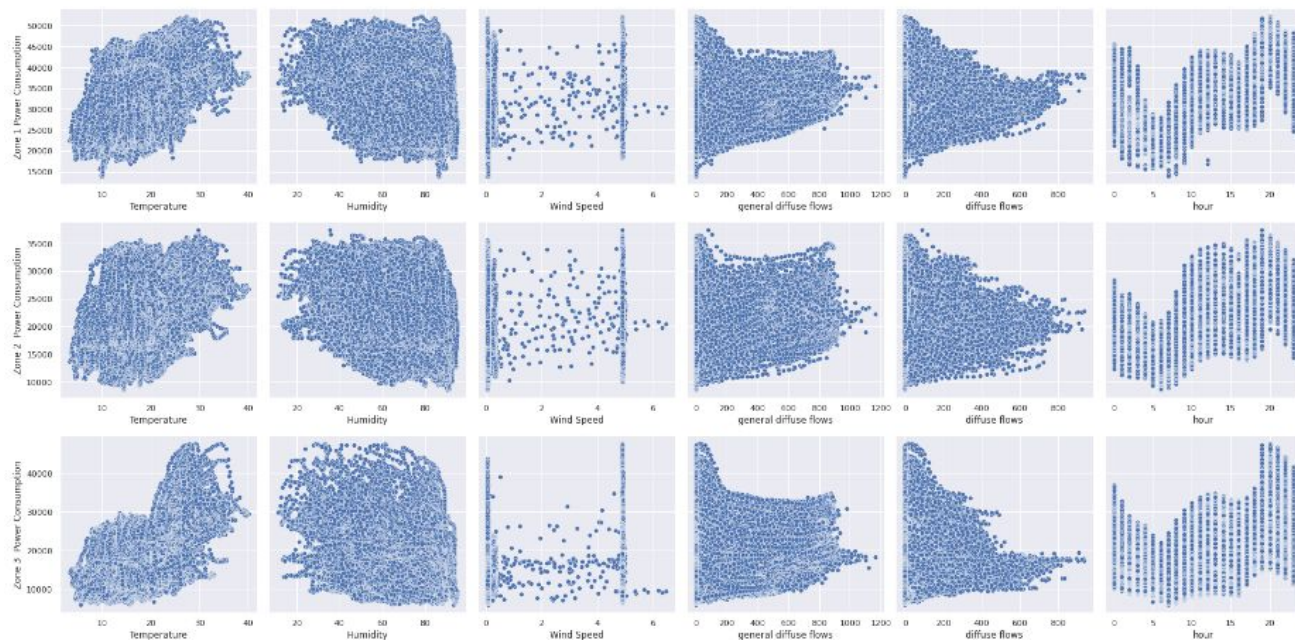
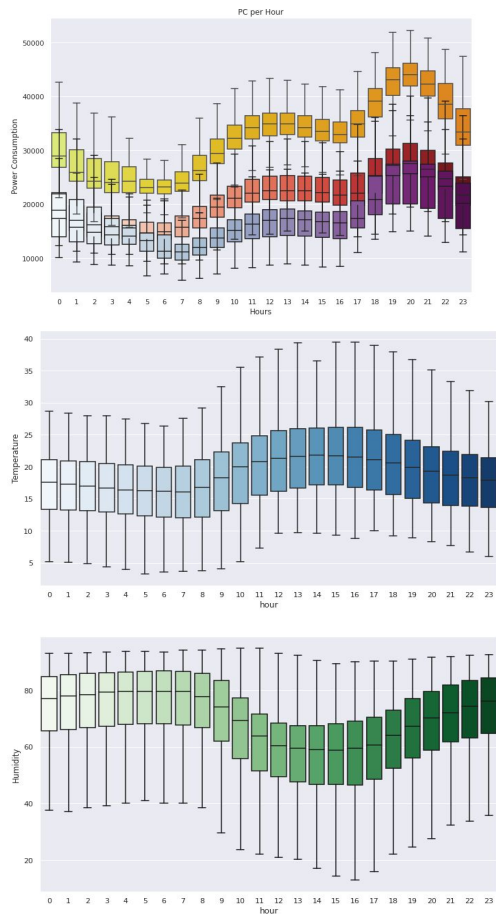
| | DateTime | Temperature | Humidity | Wind Speed | general diffuse flows | diffuse flows | Zone 1 Power Consumption | Zone 2 Power Consumption | Zone 3 Power Consumption |
|-------|------------------|-------------|----------|------------|-----------------------|---------------|--------------------------|--------------------------|--------------------------|
| 0 | 1/1/2017 0:00 | 6.559 | 73.8 | 0.083 | 0.051 | 0.119 | 34055.69620 | 16128.87538 | 20240.96386 |
| 1 | 1/1/2017 0:10 | 6.414 | 74.5 | 0.083 | 0.070 | 0.085 | 29814.68354 | 19375.07599 | 20131.08434 |
| 2 | 1/1/2017 0:20 | 6.313 | 74.5 | 0.080 | 0.062 | 0.100 | 29128.10127 | 19006.68693 | 19668.43373 |
| 3 | 1/1/2017 0:30 | 6.121 | 75.0 | 0.083 | 0.091 | 0.096 | 28228.86076 | 18361.09422 | 18899.27711 |
| 4 | 1/1/2017 0:40 | 5.921 | 75.7 | 0.081 | 0.048 | 0.085 | 27335.69620 | 17872.34043 | 18442.40964 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 52411 | 12/30/2017 23:10 | 7.010 | 72.4 | 0.080 | 0.040 | 0.096 | 31160.45627 | 26857.31820 | 14780.31212 |
| 52412 | 12/30/2017 23:20 | 6.947 | 72.6 | 0.082 | 0.051 | 0.093 | 30430.41825 | 26124.57809 | 14428.81152 |
| 52413 | 12/30/2017 23:30 | 6.900 | 72.8 | 0.086 | 0.084 | 0.074 | 29590.87452 | 25277.69254 | 13806.48259 |
| 52414 | 12/30/2017 23:40 | 6.758 | 73.0 | 0.080 | 0.066 | 0.089 | 28958.17490 | 24692.23688 | 13512.60504 |
| 52415 | 12/30/2017 23:50 | 6.580 | 74.1 | 0.081 | 0.062 | 0.111 | 28349.80989 | 24055.23167 | 13345.49820 |

52416 rows × 9 columns

Feature importance: Correlation Matrix

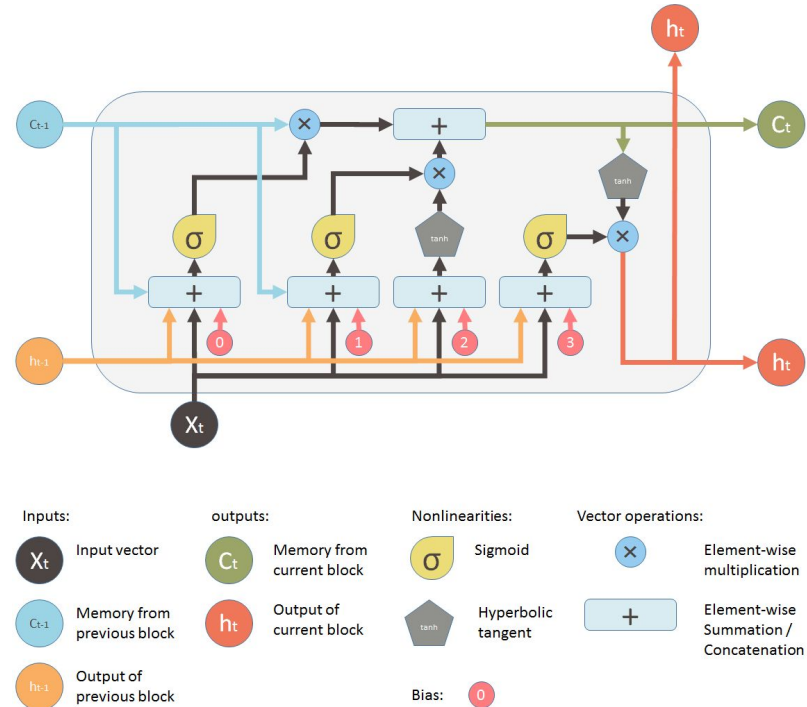


Feature importance: Data Visualization



Model #1: Stacked LSTM (Neural Network)

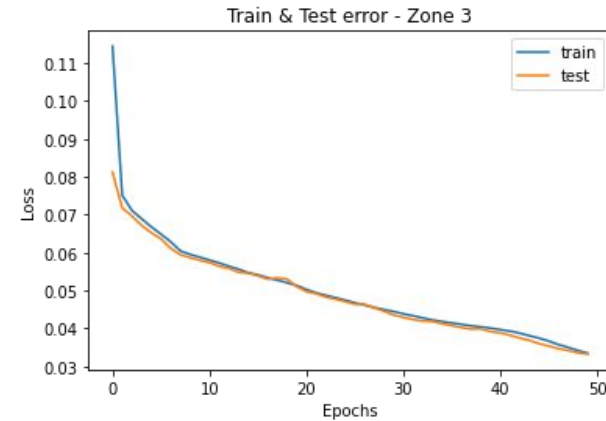


1. Neural Network structure used: Stacked LSTM
 - a. Input Layer
 - b. Hidden Layers:
 - i. **LSTM with 128 blocks**
 - ii. **LSTM with 64 blocks**
 - c. Output Layer: Dense layer with 1 neuron
2. Model Hyperparameters:
 - a. Batch Size: 256
 - b. Epochs: 50
 - c. Train-test split: 70-30
 - d. Optimizer: Adam
 - e. Loss for training and validation: MAE
 - f. Dropout: 0.2
 - g. Activation for the hidden layers: ReLU
 - h. Activation for the output layer: Sigmoid



Model #1: Stacked LSTM (Neural Network)

Accuracy measurement:

| Zone | Mean Absolute Error (MAE) | Root Mean Square Error (RMSE) | R ² |
|------|---------------------------|-------------------------------|----------------|
| 1 | 1953.011 | 2715.579 | 0.878 |
| 2 | 4206.401 | 5126.158 | 0.533 |
| 3 | 2291.988 | 3051.765 | 0.811 |






Model #2: Decision Trees

Hyperparameters used:

1. Criterion: Squared Error
2. Max Depth: 5

Feature importance of Zone 1, 2, and 3 (From left to right):



Model #2: Decision Trees

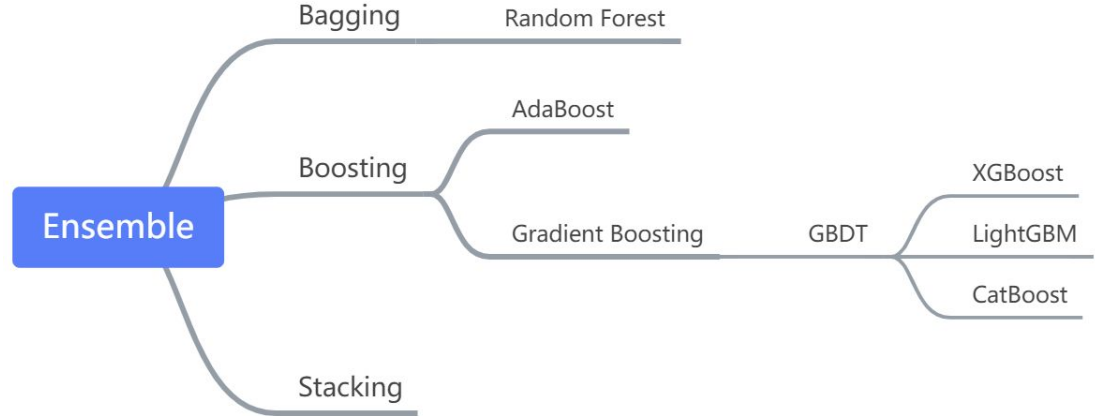
Accuracy measurement:

| Zone | MAE | RMSE | R ² |
|------|----------|----------|----------------|
| 1 | 2031.663 | 2733.539 | 0.854 |
| 2 | 1702.129 | 2246.524 | 0.812 |
| 3 | 1722.088 | 2429.634 | 0.867 |

Model #3: Random Forests & Boosting Algorithms


Ensemble methods:

1. Random Forest: Bagging
2. XGBoost: boosting(Gradient-boosting decision tree)
3. LightGBM: boosting(Gradient-boosting decision tree)
4. CatBoost: boosting(Gradient-boosting decision tree)



Difference within boosting **models**:

1. XGBoost: level-wise
2. LightGBM: leaf-wise
3. CatBoost: Symmetric trees



Model #3: Random Forests & Boosting Algorithms

Hyperparameters **tuning** (grid search)

1. **XGBoost:** learning_rate=0.015, n_estimators=1000, silent=True, nthread=1, max_depth=16
2. **Random Forest:** n_estimators = 1000, oob_score = True, n_jobs = 1, random_state = 1
3. **Light GBM:** max_depth=16, min_child_samples=0, n_estimators=4000, subsample_for_bin=2000000
4. **Catboost:** loss_function='RMSE', iterations=600, learning_rate=0.5, depth=8

Accuracy measurement (for zone-1, which can also be extended to other zones):

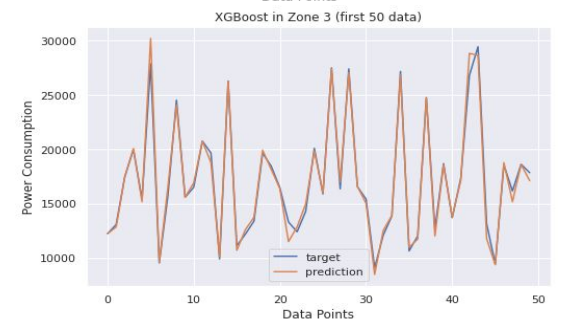




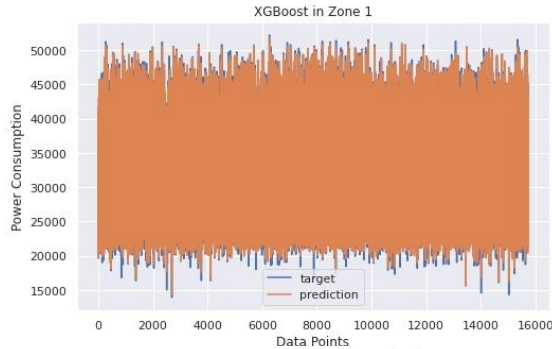
| Model | RMSE | R ² |
|---------------|----------|----------------|
| XGBoost | 950.953 | 0.982 |
| Random Forest | 1045.304 | 0.978 |
| Light GBM | 1062.366 | 0.977 |
| CatBoost | 1108.793 | 0.975 |

Summary Table

| Methods | Neural Network | Decision Tree | Ensemble Learning |
|---------------------|----------------|---------------|-------------------|
| Best R ² | 0.878 | 0.867 | 0.982 |

Hence, our best fit model is: XGBoost

Final Predictions



Conclusion

- Boosting algorithms (like XGBoost) are certainly more accurate than deep learning algorithms.
- We suspect that deep learning works better on homogenous datasets, but since ours is heterogeneous, boosting algorithms perform better as tree-based methods treat features independently of each other and build rules based on the values for those features.

Future Development

- Every Data Scientists' nightmare: Is our model overfit?
- LSTM: Implement other architectures like vanilla/bidirectional LSTM or even GRUs
- Trees: Play around with hyper parameters such as criterion and max depth for improvement
- Unsupervised learning: Try k-means algorithms



THANK YOU!

ANY QUESTIONS?

TRANSCENDING DISCIPLINES, TRANSFORMING LIVES



COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science